# DEAL CONFIGURATOR | USER GUIDE

Select the actual open-weight model, define customer demand, and understand the commercial impact of capacity choices. The four-tab workbook is an underwriting prototype, not a live AC2 pricing API.

## Start here: one input sheet

| STEP / LOCATION | DO THIS | WHAT CHANGES |
|---|---|---|
| **1 \| Deal Setup B7** | Choose a named model using the dropdown. B8 displays total parameter count (billions). | Model class, training/serving flags, scenario throughput and fee adjust. |
| **2 \| Deal Setup F22:G32** | Check active parameters (MoE), catalog evidence, training eligibility, GPU requirement and qualification state. | Unknown parameter values and unsupported training stay on HOLD. |
| **3 \| Deal Setup B9:B48** | Enter GRPO method, maximum cycles, stopping conditions, policy and rollout tokens, eval counts, measured throughput and GPU counts. | Separate training, rollout and eval GPU-hours; cycle plan and COGS. |
| **4 \| Deal Setup B79:B89 and B103** | Choose how to acquire capacity. Record spare GPUs, their dates and GPU family in rows 109-111. | Existing allocation, new PayGo/extension/reservation commitments and gaps. |
| **5 \| Deal Setup B101:B104** | Approve benchmark, weights/license, planned GPU family and SKU-specific rate only after real validation. | Changes quote status from HOLD toward READY FOR REVIEW. |
| **6 \| Economics + Offer** | Review costs, margin, availability warnings; use Offer/SOW only as a reviewed draft. | Customer-facing fee lines and evidence-based SOW remain synchronized. |

## Model choice: more than the headline B-number

**Total parameters** approximate the full weight footprint. **Active parameters** approximate what an MoE network activates per token; training gradients, parallelism, memory, long context, tool latency and accelerator kernels still matter. A 550B/55B-active model must not be priced as a 55B dense model. The workbook calculates **illustrative scenarios** using a proxy class until real benchmark data is entered and approved.

Only send a proposal when Economics Capacity shows READY FOR REVIEW, and after formal commercial/legal signoff. “READY” is not a cloud reservation or a customer contract.

# MODEL OPTIONS AND APPROVALS

The dropdown contains 30 named, family-level, research-only, serving-only or custom model options drawn from the supplied AC2 knowledge base plus selectively checked public model cards. It is **not an exhaustive current AC2 live model list**. AC training and serving catalogs may differ.

## What the model metadata means

| FIELD | SOURCE / INTERPRETATION |
|---|---|
| **Model + params** | Lookup from Guide Assumptions F:Q. Total params in Deal Setup B8; active params on the right. Missing exact facts display Not verified. |
| **Training / serving status** | Listed: explicitly named in supplied KB section. Family - verify: family mentioned but exact checkpoint not validated. Research-only does not grant platform access. Serving-only models cannot be quoted for RL. |
| **Proxy class 1 / 2 / 3** | For scenario exploration only, respectively using 2.5x / 1x / 0.5x of baseline throughput and 0.65x / 1x / 2.5x cycle fees. These are not AC measurements. |
| **GPU type / pool fit** | B103 selects proposed accelerator family; pool G109:G111 must match. B101, B102, B104 require measured workload speed, legal/access approval and SKU-specific rates. |
| **Unknown / gated variants** | Blank or unverified specs are not filled with guesses. Pick another model or obtain a verified model card and AC2 training approval. |

## Example selections

| MODEL | AUTO-PARAMETERS | ACTION |
|---|---|---|
| **Qwen3.6-35B-A3B** | 35B total / 3B active | Named training + serving; still benchmark before quoting. |
| **Kimi K3** | 2,800B total / 104B active | Large MoE; validate B300 cluster / training topology / procurement. |
| **Nemotron 3 Ultra** | 550B total / 55B active | Multi-node suitability and license check. |
| **GLM-5.3-Flash** | Not verified | Serving-listed only; BLOCK RL quote. |

Catalog maintenance: add any newly approved checkpoints to Guide Assumptions F:Q and update the B7 data-validation range (or, preferably, refresh model_catalog.json in the coded MVP). Never claim that a public model being open-weight automatically means AC2 supports RL on it.

# ASSUMPTIONS THAT MATTER

These are deliberately visible defaults. A verified model card does not validate a compute-price forecast. Replace all uncertainty-tagged values before sending a binding commercial quote.

## Model and engineering forecasts

| ASSUMPTION | DEFAULT / WHY | REPLACEMENT EVIDENCE |
|---|---|---|
| **RL program** | 3 weekly cycles; one use case; 2-week readiness and 1-week handoff. Gives a bounded SOW. | Approved customer experiment plan and stopping/plateau rules. |
| **Token and eval work** | 42.67M policy tokens/cycle; 4K rollouts; mean 12K input + 10K output; 200 evals/cycle. Demonstration workload. | Versioned customer tasks, traces, grader cost and throughput benchmarks. |
| **Model throughput** | 100K work-equivalent tokens/GPU-hour per stage, multiplied by heuristic class. Not a measured GPU benchmark. | Same model + GPU family + method + context/concurrency profiler. |
| **Labor and non-GPU costs** | $150-200/hr engineer cost, grading, sandbox, platform, storage and 10% rework. Scenario only. | Internal compensation burden, signed staffing plan and historical vendor invoices. |
| **GPU reservation** | 8 prepaid H100 GPUs weeks 1-8, allocated at $2.50/GPUh; PayGo $3.99/GPUh. Explains existing-vs-new economics. | Real dated inventory, region, interconnect, supplier contract and quoted SKU price. |
| **Pricing + margin** | $30K readiness; $78K integration; $24K/standard cycle; $30K handoff; 45% target GM. Not company policy. | Approved rate card, finance margin policy and customer willingness-to-pay. |
| **Production** | Optional separate order; 2 warm GPUs and illustrative token traffic. Shows idle-capacity floor. | Per-model SLO load test, p95/p99, traffic peaks and failover design. |

## Important limits

Planning covers only 16 weeks and three capacity tranches, and it does not prove physical cluster topology, supply availability, model licensing, GPU quota, performance SLO, or a genuine time-to-completion guarantee. Customer retains stop/continue decision; cycle overages, new base models and production expansion require written commercial review. See the full audit for material fixes and remaining limitations.