MODEL / ECONOMICS AUDIT

Review scope: four-tab deal configurator, prior guide, MVP brief, sample configuration, model claims in the supplied 2026-10-18 Applied Compute knowledge base and external model-card checks.

Issues corrected in the workbook

RISK IDENTIFIED	CORRECTIVE ACTION
Model tier disconnected from exact model	Exact model dropdown and auto total+active parameter lookup. Selected model now drives planning class, GPU requirements and cycle price.
Dense-size proxy could misprice MoE frontier model	Models now display architecture, total and active parameters; benchmark approval is required before commercial release. Large models remain provisional.
Unsupported model could be quoted	Training status distinguished from serving; unlisted, family-only, research-only, gated and missing-parameter selections cannot be approved.
Pool “Compatible” did not identify H100/B300	Added GPU family to each existing tranche and planned GPU type; type mismatch removes supply from eligibility.
Zero throughput/invalid workload broke arithmetic	Zero-safe formulas plus HOLD - INPUT prevent invalid cost assumptions from becoming approved quotes.
Customer proposal exposed premature price	Customer Offer/SOW now explicitly says DO NOT SEND while HOLD; any proposed margin-floor adjustment needs customer approval.
Serving eligibility treated as implicit	Optional serving now checks whether selected checkpoint is serving-listed and separately benchmarked.

Regression tests completed

Baseline (35B/3B), 9B family case, 2.8T Kimi B300 mismatch, missing-spec serving-only model, approved model, fourth RL cycle, GPU SKU changes, existing-vs-new capacity, and formula error scan. All propagated and showed appropriate HOLD/READY states. This tests formula mechanics—not supplier prices or production performance.

UNVERIFIED ASSUMPTIONS

Material unknowns do not all “break” a formula, but many prevent using the result in a real contract. The spreadsheet keeps them visible and enforces mandatory approval gates.

Critical before quoting

AREA	CURRENT BASIS	VALIDATE BEFORE RELEASE
Actual AC2 model availability	30-item named/family/research snapshot; live model catalogs were not fully retrievable.	Current training-vs-serving endpoint status, SKU eligibility, access control, checkpoint/version and applicable license.
Model sizes and architecture	Some first-party cards verified; some secondary; some absent. Active count may be direction-dependent.	Original model weights/config; quantization and input/output active parameters as applicable.
Performance and memory	Proxy class 2.5x/1x/0.5x, 8/8/4 base GPUs.	Measured GRPO policy, rollout, grader and inference throughput, GPU memory, topology and end-to-end experiment wall time.
Capacity + cost	8 spare H100s; provider availability unconfirmed; one public H100 PayGo benchmark.	Dated reservation inventory, node layout, supplier confirmation, B200/B300 rate, contract expiry and other customer allocations.
SOW + economics	Engineering rates, delivery uncertainty, 45% target margin and rate card are illustrative.	Finance-approved allocations, real delivery hours, technical customer signoff and commercial policy.
Serving readiness	Warm replica and token-equivalent forecast omits engine-specific latency/HA.	SLO benchmark, input/output work split, HA/failover, rate limits, peak demands, operational coverage.

Core external references

AC2 training: https://docs.appliedcompute.com/platform/training/models
AC2 serving: https://docs.appliedcompute.com/platform/serving/models
Qwen3.6: https://huggingface.co/Qwen/Qwen3.6-35B-A3B
Kimi K3: https://arxiv.org/abs/2607.24653
Nemotron Ultra: https://github.com/NVIDIA-NeMo/Nemotron
DeepSeek V4.1 Flash: https://deepseek.com/en/news/deepseek-v4-1-flash/
GPU proxy: https://lambda.ai/pricing

When in doubt: an approved customer price requires verified cost and capacity, not merely a plausible model size.