CUSTOMER PRICING / ORDER FORM

Enterprise post-training quote

One workflow, one selected base model, three customer-approved RL cycles. This is an **illustrative offer** for a proposed six-week engagement, pending technical fit and final approval.

A. Included services and fees

Item / description	Quantity	Unit price	Extended
Readiness + baseline. Data assessment, model recommendation, grader and held-out baseline.	1 phase	$30,000	$30,000
Implementation engineering. Harness integration, graders, run setup and monitoring.	300 combined hours	$260/hour	$78,000
RL improvement cycles. Three supervised weekly cycles incl. run/eval and compute allowance.	3 cycles	$24,000/cycle	$72,000
Final validation + handoff. Blinded evaluation, comparison, checkpoints, configs and report.	1 phase	$30,000	$30,000
TOTAL SIX-WEEK ENGAGEMENT			**$210,000**

Included operating envelope

Scope: one enterprise workflow, one approved model family, three RL cycles (target cadence: one per week after readiness), baseline plus final evaluation, and 300 implementation engineering hours. The readiness and closeout work are fixed service packages separate from those 300 hours.

Compute allowance: up to 3,360 charged H100-equivalent GPU-hours across the three-cycle package (planning average 1,120 per cycle), assuming the approved standard model and benchmarked GPU topology. The allowance is **included in the cycle price**, not separately charged.

Approval: the customer may stop after any cycle. Work beyond the authorized quantity is performed only under written change approval. Unused cycles are not automatically refundable; phase-level cancellation terms are negotiated in the order.

Payment schedule

Billing event	Amount
Phase 0 kickoff	$30,000
Phase 1 authorization	$90,000
Midpoint after agreed cycle evidence	$54,000
Final handoff and acceptance	$36,000

Prices and timing are illustrative; any transaction taxes, externally required security work, materially different architecture, or new base model require a revised order. See the standalone SOW for acceptance and decision rights.

CUSTOMER PRICING / OPTIONS

Model choices and change rates

Model size affects throughput, GPU topology and required experimentation. These are **planning tiers**; final eligibility and price follow a benchmark of the selected weights, context length, training method and harness.

B. RL cycle package by model class

Illustrative model class	Price per extra cycle	Included compute / cycle
Small: 4–8B dense	$12,000	up to 560 H100-equivalent GPU-hours
Standard: about 20–40B active/dense	$24,000	up to 1,120 GPU-hours
Large / sparse-MoE / multi-node	$60,000	up to 2,800 GPU-hours; subject to feasibility

Each cycle includes an approved starting checkpoint, one agreed workload run, fixed evaluation set, cost/quality report, and customer stop-or-continue gate. GPU-hour allowances are **limits, not guaranteed performance**; work is constrained by approved workload and provider availability. A model change triggers re-benchmarking and a revised commercial order.

C. Additional work: only with written approval

Charge type	Illustrative unit price	When billed
Additional RL cycle	$12K / $24K / $60K	Selected model-class package; not an automatic renewal
Forward-deployed engineer	$280 / hour	Additional workflow / harness integration
ML research engineer	$320 / hour	Additional reward, training or model analysis
Infrastructure engineer	$260 / hour	Nonstandard cluster / deployment integration
Project or delivery lead	$200 / hour	Extra workshops or documentation
Excess approved GPU consumption	$9 / H100-equivalent GPU-hour	Only above included allowances with written order; no duplicate billing for included hours

Cycle cadence and decision point

After baseline approval, one improvement cycle is scheduled in each of weeks 3, 4, and 5. Applied Compute submits a short run/eval report after each cycle. The named customer decision-maker approves the next cycle, elects to stop, or requests a scoped revision.

Not included: an additional workflow, a new model family, production uptime commitments, a new region, continuous RL, indefinitely repeated runs, customer-required external data licensing, or services beyond the agreed deliverables.

CUSTOMER PRICING / INFERENCE

Optional inference services

Inference evaluation and production are **separate orders**. The prices below are discussion examples and are not included in the post-training engagement above.

D. Evaluation pilot and production

Product	Proposed charge	Included service / boundary
Two-week inference pilot	$12,000 / pilot	One model; test harness, capped replay of up to 2,000 requests; model-specific token usage allowance stated in pilot order
Dedicated production serving	$18,000 / month, plus metered tokens	Two warm replicas, single region, one agreed model and measured service envelope; subject to load test
Extra warm replica	$7,000 / replica-month	Requires capacity and service-order amendment before deployment

E. Illustrative token rate menu (separate from monthly capacity fee)

Open-weight deployment	Input / 1M tokens	Output / 1M tokens	Rate basis
Small: 4–8B dense	$0.35	$1.00	Hypothetical planning rate
GLM 5.2	$1.40	$4.40	AC public calculator example*
Large / multi-node	$3.50	$10.00	Hypothetical; quote after test

*Applied Compute publicly illustrates GLM 5.2 at $1.40 per million input and $4.40 per million output tokens. Those rates are model-specific examples, not guaranteed rates for other models or this offer. Source: https://www.appliedcompute.com/inference.

F. Initial production envelope and scaling rules

The working production order would cover up to **two warm replicas, 20 concurrent requests, and 100 million aggregate tokens per month**, only after a representative traffic replay establishes achievable throughput and latency. Token charges are in addition to the stated capacity fee. A capacity review begins at 70% of the agreed service-level-compliant throughput for two consecutive weeks, or before forecast volume/concurrency exceeds the order.

Additional replicas, more demanding latency targets, higher uptime, a different region or model, or a longer reserved-GPU term require an agreed service amendment before capacity is procured. Production SLA and credit terms must be included in that later service order.

Terms of quote

Customer acceptance and required approvals are defined in the SOW. No automatic additional research or infrastructure spend. Model choice, workload assumptions, token rates, service envelopes, and availability are subject to written validation. All fees above are illustrative, not official Applied Compute price disclosures.