CUSTOMER SOW / DELIVERY

Model improvement: scope & evidence

One customer-approved workflow and base model, from readiness to three reviewable RL cycles and final handoff. This document defines **work and acceptance**; the separate order form states all charges.

Phased delivery and ownership

Phase / timing	Applied Compute delivers	Customer owns	Evidence to accept
0. Readiness Weeks 1–2	Data + harness review; model alternatives; signed measurement plan; baseline run.	Provide lawful data access; approve base model, evaluator and held-out set.	Dataset / grader version hashes; baseline scores; test definitions; go/no-go record.
1. Setup Week 2	Implement training environment, graders, run logging and resource plan.	Authorize security/compute profile; validate grader against SME judgment.	Harness smoke test; reproducible scorer; approved cluster and run manifest.
2. RL cycles Weeks 3–5	Run one cycle weekly; capture checkpoints, run costs, deltas and failure analysis.	At each gate approve continue / stop / change within remaining scope.	For EACH cycle: run ID, config, tokens/GPUs, trace sample, quality deltas, regression notes, decision log.
3. Final evaluation Week 6	Run held-out evaluation, provide weights, configs, performance/limitations report and transfer.	Run acceptance review; confirm operational handoff and IP provisions.	Final checkpoint manifest, blinded evaluation vs baseline, grader output, reproducibility steps, signed acceptance.

Evidence standards (not an unconditional improvement warranty)

Baseline: signed data split, labeled cases, success metric, scorer version, reproducibility seed and initial result. **Per cycle:** a traceable run and checkpoint, consumption summary, quality and safety regression checks, and a stop/continue recommendation. **Final:** held-out comparison, results by task cohort, material failure modes, checkpoint export and the steps to rerun the tests.

Acceptance occurs when the agreed outputs and evidence are provided and meet documented test procedures; a minimum quality uplift is warranted only if expressly stated after readiness. Customer provides a written acceptance or a precise list of unmet SOW criteria within five business days; disputes proceed through a remediation review, not automatic acceptance.

CUSTOMER SOW / CONTROL & EXPANSION

Decision rights and when scope resets

Applied Compute recommends and executes technical work; the customer defines success, approves spending decisions, and determines whether additional training has business value.

Decision / responsibility matrix

Decision / action	Applied Compute	Customer
Initial base model and licenses	Recommend feasible weights, architecture and model trade-offs.	Select/approve initial base model, intended use and license obligations.
Data, grader and held-out test	Build, version, validate, and flag grading risk.	Own data rights and subject-matter rubric; sign off success definitions.
Weekly research cycle	Execute approved cycle, flag regressions, report measured consumption.	Explicitly approve start of each next cycle or stop remaining research.
Acceptance and model ownership	Deliver checkpoint, code/configs, grader output and transfer records.	Accept or reject against evidence; retain customer rights per governing IP agreement.
Change in work or deployment	Assess incremental engineering, security, GPU and timeline impact.	Authorize a new SOW/order BEFORE additional work or reserved capacity.

Production serving and the initial scaling boundary

Production is **excluded** from this training SOW. Following a separate approved inference pilot and service order, a planning envelope could include **one region, one model, two warm replicas, up to 20 concurrent requests and 100 million tokens per month**. These are proposed commercial limits, not guaranteed hardware throughput until the load test establishes latency and throughput targets.

Applied Compute would monitor traffic, request latency, throughput, replica saturation, GPU headroom and forecast demand; recommend a scaling change at **70% of contracted sustainable capacity for two consecutive weeks**, or earlier if the forecast approaches concurrency or token thresholds. Scaling past the signed envelope is **not automatic** and requires the customer-approved revised service order.

Amendment or new service order: precise triggers

Required for another initial/base model or materially different weights or license; another workflow or grader; further RL cycles or workload beyond the included envelope; new dedicated GPUs or a supplier term extending past the customer commitment; production serving/SLA; added replicas or higher concurrency/volume; new region or data isolation tier; or continuous retraining. AC provides a refreshed benchmark, capacity requirement, delivery timetable and revised price before acceptance.

Contract controls across the customer lifecycle

During RL: scope is limited by workflow, base model, scheduled cycles, engineering and GPU allowances. Customer approval is required before extra research or supplier commitments; cancellation treatment for unavoidable supplier reservations is defined by the commercial order. **At serving scale:** the separate order must align minimums, reserved-capacity term, uptime, burst limits, renewal and early termination with actual infrastructure exposure.